CLIP-based deepfake detection — corrected v5 dashboard (E14 strict LODO)
Four protocols, n = 1758 subject-aware throughout

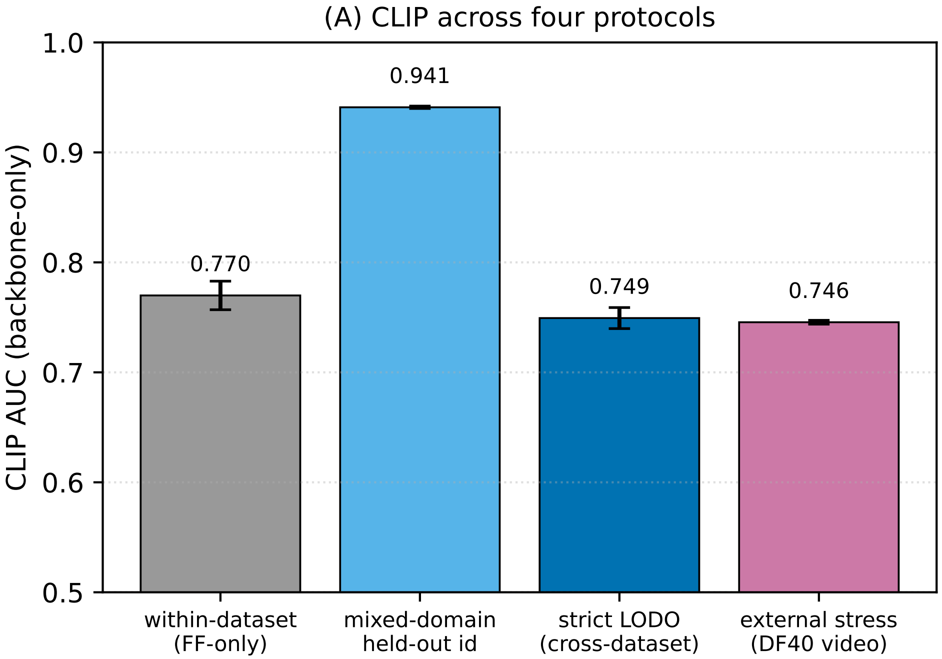

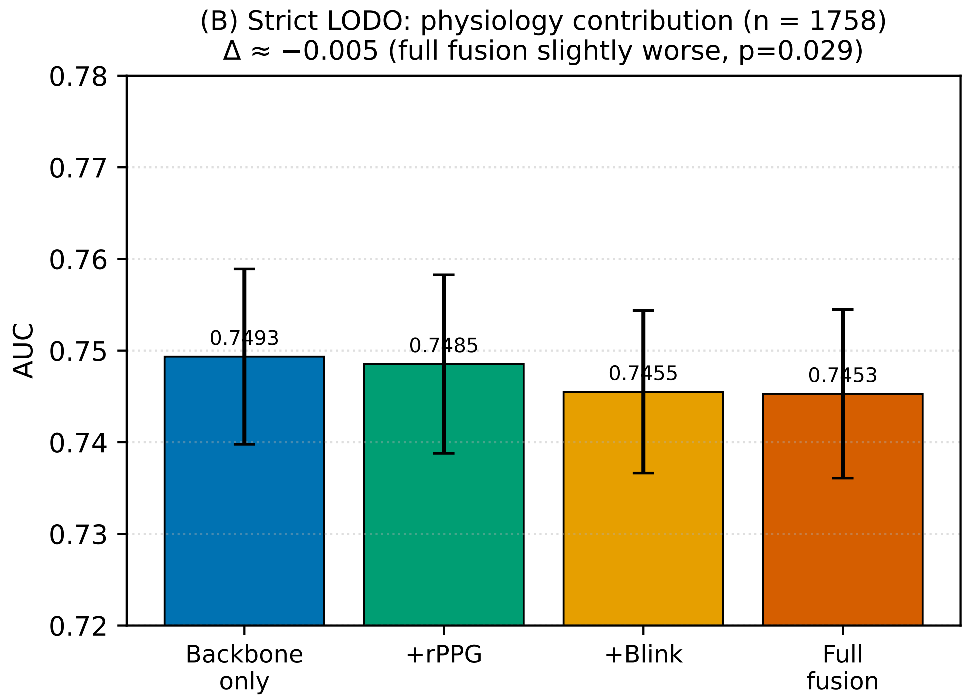

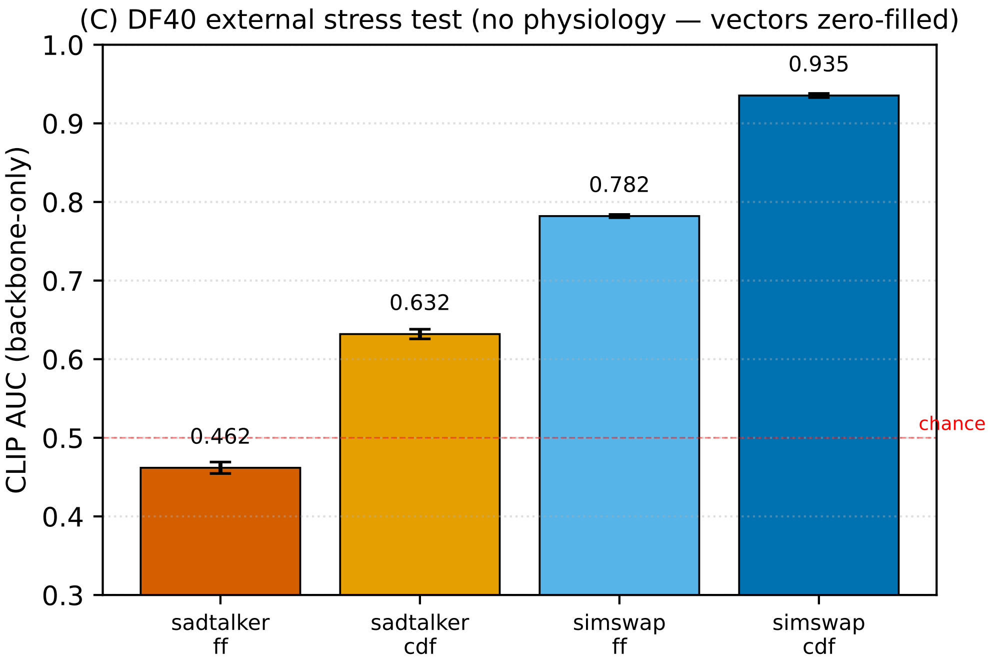

Headline findings (v5 − corrected, E14 strict LODO)

1. CLIP CelebDF AUC = **0.749 ± 0.010** under strict LODO (n = 1758, train FF + DFDC, CelebDF fully held out). v3/v4 0.941 number was a mixed-domain leak — retracted.

2. Full fusion = 0.745; Δ = −0.005 vs backbone-only. Stouffer p = 0.029 in favour of backbone-only — physiology is at best neutral and slightly negative under strict LODO.

3. Apples-to-apples vs DeepfakeBench (FF-only → CelebDF): CLIP 0.770 ≈ SPSL 0.7650 (field ceiling) without auxiliary hand-engineered features. Strict LODO 0.749 ≈ Xception 0.7365.

4. Physiology has *representation−dependent marginal value*: small residual gains for B4 within-dataset, neutral or slightly harmful for CLIP cross-dataset.

5. DF40 stress test: CLIP 0.746 ALL, but sadtalker_ff 0.462 (below chance) — talking-head reanimation fundamentally harder than face-swap.